

CSC 447/557 – Introduction to Computer Vision

Assignment One

Due: 11:59pm (*) Thursday, January 23.

(*) There is grace until 8am the next morning, as the instructor will not grade assignment before then. However, once the instructor starts grading assignments, no more assignments will be accepted.

Weight about 5 points

This assignment should be done individually

The purpose of this assignment is to become familiar with some basic image manipulation in Matlab. Matlab is useful for exploring ideas and prototyping programs. It is a popular programming environment for computer vision and machine learning, especially for those who do not have lots of experience in C/C++ because it allows one to focus more on the problem and less on the code. It has significant drawbacks for writing large programs or when performance matters. It is also a “product” which needs to be purchased and installed before it can be used.

Nonetheless, it is a useful tool which computer vision and machine learning students should be exposed to. It is likely easiest to do this course in Matlab, but you can use whatever language you like for subsequent assignments. To the extent that I provide any coding solutions at all, they will be in Matlab.

Matlab is available for personal use to UA faculty, staff and students for free. It can be downloaded from <http://sitelicense.arizona.edu/mathworks-matlab/> by any University of Arizona faculty, staff or student with a netid. The site-license includes MATLAB, Simulink, plus 48 toolboxes. It is also installed in a number of UA machines and labs. If you do not have convenient access to Matlab, please contact the instructor to discuss options. **Make sure you download the image processing toolbox.** If Matlab does not know about the command "imshow", then you probably do not have this toolbox. Note that if you download all of Matlab (instead of picking and choosing your toolboxes), it can take a long time, especially on a slow connection. Be prepared for that!

Deliverables

Deliverables are specified below in more detail. For the high-level perspective, you are to provide a program to output a few numbers and to create figures based on the input image Data/tent.jpg posted on D2L. (D2L might not show the ".jpg", depending on settings). There is no need to upload this image to D2L when you submit your homework.

You also need to create a PDF document that tells the story of the assignment, including output, code snippets, and images that are displayed when the program runs. Even if the question does not explicitly remind you to put the resulting image into the PDF, if it is flagged with (\$), you should do so. The TA should not need to run the program to verify that you attempted the question. See the file Info/assignment-instructions.pdf under D2L for more details about doing a good write-up. You are also encouraged to have a look at the more thorough "Info/writing-tips.pdf". While writing good reports takes work, it is well worth getting better (and more efficient) at this. A substantive part of each assignment grade is reserved for exposition. In this assignment a number of parts do not have “deliverables”; they are just to help you learn your tools. These are flagged with “@”. The parts of the assignment that imply a deliverable are indicated with (\$). Note that each deliverable figure created as the program runs should be a new one. If (Matlab) you do not use the “figure” command as explained below, a new figure overwrites the previous one.

1. Prepare for next week (@, no deliverables).

Next week we will use some very basic linear algebra. Please review, if needed: computing dot products, matrix vector multiplication, and matrix multiplication. Also, consider what it means to approximate a simple function using a vector. Finally, do you remember what it means, formally, to be a linear function?

2. Become familiar with Matlab (@, no deliverables).

Spend some quality time with a Matlab tutorial (there are many on-line), or perhaps just look at the introductory material that is part of Matlab's extensive help system. To do the later, start up Matlab and hunt for a "Help button" (on Matlab R2016a on my mac there is one in the menu bar and one towards the right hand side of home top bar). From there, you probably need to choose "Documentation". Then you need choose help for Matlab itself (instead of one of the tool boxes and other accompanying products). From there, I suggest exploring "Getting started" and "Language Fundamentals".

The complete documentation for Matlab is also available on online here:

<http://www.mathworks.com/access/helpdesk/help/helpdesk.shtml>

but you might notice that these pages are the same as the ones you look at using the Matlab desktop help browsing system (which might explain why you might find the help system bit clunky --- the information is arranged in a page tree instead of a true narrative).

Regardless of how you do it, it is a good investment to go through some introductory material first to get an overview of how it works. It will save time in the long run.

Documentation for Matlab commands is also available from the command line, which is especially useful to remind yourself about the details of a function you know the name of, or a function mentioned in this assignment. For example, if you want to know how the `svd()` function works, type `help svd`. You can also get help on all the built in operators and even the language itself. Typing `help` gives a list of help topics.

*Tip: You may want to turn on paging (**more on**) for reading help pages.*

3. Reading and Displaying Images (@, no deliverables)

Place the JPEG color image `Data/tent.jpg` posted on D2L (also linked here: <http://kobus.ca/teaching/cs477/data/tent.jpg>) in your working directory.

Load the image into a variable using `imread`. Type `help imread` to find out how to do this. Create a new figure using `figure` and display the image in it using `imshow`.

Tip: Make sure you put a semicolon at the end of your `imread` command. The semicolon at the end of a line prevents Matlab from displaying the result of an expression. Most of the time, you want the semicolon.

4. Writing Images (@, no specific deliverables, but you will need to do **something** to get images into your PDF report!)

Write the image to a file called `out.jpg` using the `imwrite()` function. Use some independent image viewer such as `display`, `xv`, or a web browser to verify that this worked. If you are a Linux user (Mac and Windows also, but even more likely you need to install), I recommend learning about the ImageMagick suite of tools for converting, and displaying images (do "man convert", and a "man display" to find out more). Also the program `import()` can be used to get a screen-shot (at least in Linux; "grab" works best on Mac, and I have no idea about Windows).

5. Scripts, files and paths (@, no specific deliverables, but you need to follow these instructions to complete the assignment)

The interactivity of Matlab is great for debugging and experimenting, but often one wants to type code into a file. As you develop code for the deliverables, put them into a file called `hw1.m`. You can then type `hw1` at the Matlab prompt to execute the commands in the file. A file of this sort is known as a script.

Tip: You might prefer the output if you use `format compact`.

Tip: My scripts often have `close all` in the beginning, so when I re-run them the old images being displayed go away. Please do this for assignments in this class to make the TA's life easier.

Tip: Matlab has `pwd`, `ls`, and `cd` commands that do what you expect.

Tip: If your script is in some other directory than `pwd`, then you can add that other directory to Matlab's search path with the `addpath` command. The current directory is in Matlab's search path by default.

6. Functions (@, no specific deliverables, but you need to follow these instructions to complete the assignment)

You can define a new Matlab function by simply putting code in a file (as we have just done) and placing a function declaration at the top. Write the following as the first line of your `hw1.m` file:

```
function [num_questions] = hw1(infile)
```

Tip: The name of the dot-m file (without the dot-m) and the name of the function should always match. A file can export only one function, but the file may contain internal helper functions. If you use such functions (likely you will want to in future assignments), you will need to end them with “end”---see “help function” for details. Matlab will read the file to find such internal functions before it starts executing instructions within the file.

Modify your code so that the `imread` command use the file provided by the `infile` variables instead of the hard-coded “tent.jpg”. Also, make it so that every question that you implement (starting with the next one) adds one to a variable `num_questions` so that this number can be returned by the function. The TA should be able to execute this function by typing `hw1('tent.jpg')`, or `n=hw1('tent.jpg')` if they want the number of questions done returned in a variable (The TA will check this).

*Tip. Notice that Matlab strings are specified using **single** quotes.*

Tip: If you want to return more than one variable from a function, add it to the list in square brackets in the function's declaration. Simply setting these variables inside the function will cause them to be returned.

Tip: As is common with many languages, functions do not change the values of input parameters (call by value). They can be changed (as copies) within the function, but the version of the variables outside do not change. To get values out, use the output parameters. (For pedants: there are omitted details here that are not relevant to getting started).

Tip: Functions can also have optional input and output arguments. To determine the actual number of input and output arguments a function was called with, look at the special variables `nargin` and `nargout`.

Problems that have explicit deliverables with respect to your PDF report start here.

7. Documenting Functions

When you create a new exported function, it should always be documented so that `help` returns something informative. The convention in Matlab is to place the help message in comments after the function declaration (the comment character is “%”).

For example, you can look at some of the code in the Matlab library. Some of it is implemented as .m files, and some of it is “built-in”. Regardless, there should be a .m file for at least the documentation. Using the `which` command, see if you can find the .m file for your favorite function so far. On my mac, the path provided was not exact, but gave me a good idea where to look.

If you look at some of these examples, you will see that the first line of the comment contains a one-line description of the command. All subsequent contiguous comment lines are included in the help message. Add a few lines of text about your program in this style, and verify that typing “`help hw1`” does what you expect. Update the description of what your program does as you work on it. Finally, when you are done, and paste the result of “`help hw1`” into your PDF report (\$) as the deliverable for this (first) question with deliverables.

8. Basic image data structures

Make sure you understand the data structure being used to represent the image that you have read in. Type `whos` to get a list of active variables along with their types and sizes. You can see that your image is a 3D array of bytes. This is row by column by “channel” where the channels are red, green, and blue intensity values. Now use the `whos ()` function to show the information for only the variable containing the image (\$).

The function `size ()` is used to get array dimensions of an array for further processing. Show off your ability to read help files to put the number of rows, columns, and channels, into variables `num_rows`, `num_cols`, and `num_channels` respectively, using only one statement (\$).

What is the range of values contained in the image array? Use the `min` and `max` functions to output this information for red, green, blue and overall, and report it into your writeup (\$).

*Hint: These commands by default operate on only **one** dimension of their argument. You can convert your image into a 1D arrays (a vector) using the `(:)` notation before you pass it to `min` and `max`. There are many other ways to do this. I encourage you to find at least one other way.*

Hint for the hint: `help colon`.

Convert the image to grayscale using `rgb2gray`. Check the representation of the image now. Use `size` or `whos` to find this out. You will see that we now have a 2D array, or a matrix. Create a new figure and display the black and white image (\$).

9. Image channels

Create three grayscale images based on our color image by extracting the red, green, and blue pixels, respectively, into an array. Display the three images (\$,\$,\$). Are they what you expect? Can you explain why some areas that are bright in the color image are dark in some of the grayscale “slices”? (\$). Don't forget that while your program needs to display images, they also need to be part of your narrative in your PDF, and some commentary about them should there as well.

Tip: While Matlab has C style loops, you should try to avoid doing this sort of thing with low level loops. Instead, use the colon notion (`help colon`).

Create a new color image that has the green channel from the original image as its red channel, the blue channel from the original as its green, and the red channel from the original as its blue. Create a new figure and display the result and comment on it (\$).

Hint: Again, you might find the colon notation helpful (`help colon`).

10. Manipulating Matrices

For this question, we will work with the converted grayscale image created in problem 8 with double values in the range [0,1].

Hint: Use the `double` function to do the type conversion, and then divide by the maximum allowable value for a byte, which is 255.

Tip: Array indices in Matlab start at 1, not at 0 like C.

Tip: As in standard mathematical notation, the first index of a matrix is the row, and the second index the column. When viewing a matrix as an image, this means that the first index is the y direction going down, and the second the x direction going to the right. As is common with image manipulation tools, the origin is at the top left corner, the positive x axis points right, and the positive y axis points down.

Use nested `for` loops to set every 5th pixel of every 5th row to 1. This should set 1/25th of the pixels to white in a square lattice pattern. Create another figure and display this result in it with `imagesc()`(\$), and then do the same with `imshow()`(\$). Comment in your document about the appropriateness of the one function over the other for this use (\$).

Hint: Use the colon operator to define the limits of the `for` loops. See the help for `colon` and `for` to see how to do this. Specifically, you want the `minval:interval:maxval` form.

Grad students (and ugrads looking for extra problems): It might be more natural to do question 16 next. Some students prefer that I organize questions in co-convened classes by what the two groups are responsible for, but this means that grad student problems that extend undergraduate ones can appear after other problems. I will try to provide pointers like this one. Feedback on this (or anything else) is always welcome.

11. Histograms

A histogram divides up your data space into boxes, and puts counts of occurrences into them. They are visualizations of the empirical probability distribution of your data (e.g., how likely is it to come across a very red pixel?). If you are not familiar with histograms, you should read the Wikipedia article and/or some other resource about them as we will assume that you know what they are.

Convert each of the three color channel “slice” images into 1D vectors and provide histograms for them with 20 bins for each of them (\$,\$,\$) using the `histogram()` function. Again, remember to put this in your document with a caption explaining how the figures are a (limited) representation of the image).

Tip: You need to give `histogram()` 1D vectors, and if you give it something else (like a matrix), it does something different. Make sure you give it 1D vectors. This can be done using colons.

12. Plotting

Explore the `plot()` command. Plot the `sin` function over the domain `-pi:pi`. Use the `linspace` command to define the domain `x` and then do `plot(x,sin(x))`. Use the `hold on` command to plot another function on the same graph. Do this to add `cos()`. Use a different color. The running of `hw1.m` should produce a plot along these lines (\$).

13. Playing with Linear Algebra

Matlab is a great tool to for experimenting with linear algebra. The next few questions are about doing so.

Use the fact that `inv()` inverts a matrix to solve for $X=(x,y,z)$:

$$\begin{array}{rrcr} 3*x & + & 4*y & + & z & = & 9 \\ 2*x & - & y & + & 2*z & = & 8 \\ x & + & y & - & z & = & 0 \end{array}$$

Verify that your “solution” works. Make sure that your program outputs the answer, and also the “proof” that it is correct (\$).

While you should know how to invert a matrix using Matlab, using matrix inversion to solve equations is not the best way. Matlab provides a faster, and potentially more robust, method through the function `linsolve()`. Check that you get a similar (but not necessarily **exactly** the same answer using that function, and that you can invoke `linsolve()` implicitly with the “\” operator (\$). To see if `linsolve()` gives exactly the same answer, subtract the two and report the result (\$). Note that this is not the same as simply observing that they are the same when reported with a few decimal places of accuracy. If it is exactly the same, then the difference should be zero. If it is not, can you explain why some difference is reasonable? (\$).

The next two questions are optional and should not be handed in. We will encounter the methods in future assignments, and cover them in class at that time. However, if things have been going well up to this point, you may want to do these questions as a warm up.

14. Playing with Linear Algebra II (@)

If there are more equations than unknowns, then, in the general case, “classically” the problem is over constrained and there is no solution. However, in this course, we will often be assuming that such equations are approximations and have errors due to noise or other reasons, and that an exact solution cannot be found regardless. Thus, we will want to find the “best” solution. This is known as solving the equations in the least squares sense. The solution for $AX=b$, where A has more rows than columns, is given by $X=inv(A'*A)*A'*b$, where $inv(A'*A)*A'$ is known as the Moore-Penrose inverse of A . (Note that the single quote in this context takes the transpose of what is before it). Use this to solve for $X=(x,y,z)'$ in:

$$\begin{array}{rrcr} 3.0*x & + & 4.0*y & + & 1.0*z & = & 9 \\ 3.1*x & + & 2.9*y & + & 0.9*z & = & 9 \\ 2.0*x & - & 1.0*y & + & 2.0*z & = & 8 \\ 2.1*x & - & 1.1*y & + & 2.0*z & = & 8 \\ 1.0*x & + & 1.0*y & - & 1.0*z & = & 0 \\ 1.1*x & + & 1.0*y & - & 0.9*z & = & 0 \end{array}$$

Your program should output the solution and the magnitude of the error vector, i.e., $\|AX - b\|$.

15. Playing with Linear Algebra III (@)

Recall that an eigenvector of a matrix A is a vector v , so that $Av=k v$, for some scalar constant k . If A is real and symmetric, then A has real eigenvalues and eigenvectors. Note that for a random matrix, $R, R*R'$ is symmetric. (Try it!). The Matlab function `eig()` gives you eigenvectors and eigenvalues. Use these hints to create a 4x4 matrix A , and a corresponding vector v , that satisfies the eigenvector equation $Av=k v$. Show that your A and v have this relation by printing out the value of $A*v./v$.

If you are an undergraduate student, and you have done all the questions until here, then congratulations, you are done! However, if you are looking for alternatives, or extra problems, feel free to keep reading.

16. Matrix manipulations without explicit loops (*). (Remember that problems with an asterisk are required for grad students only)

This problem is related to problem 10, and is recommended for undergraduates who only want to do one extra problem. For this problem, we will use the same converted grayscale image with double values in the range [0,1] that we used in problem 10. Similar to problem 10, set every 5th pixel, horizontally and vertically, to 0 (instead of 1), so that 1/25th of the pixels are black in a square lattice pattern. However, do it this time without using any for loops in the code. This can be done in one, relatively short, statement. Create yet another figure and display this result in it (\$). Provide a code snippet in the PDF as well.

Hint: You can index arrays in Matlab with vectors as well as with scalars, so `im(Y,X)=0` will set multiple entries of `im` to zero when either `X` or `Y` are vectors, such as the vectors returned by the colon operator.

Hint on hint: What is described in the hint might take some getting used to. Like many things in Matlab, it is easiest to combine reading of the documentation with experimentation. It is helpful to try simple things on random matrices. To get a random square matrix of size 10x10 use `rand(10)`, and to get one of size 4x12 use `rand(4,12)`. A specific vector can be created by `v=[2 4 6 8]` (and a matrix by `m=[1 2; 3, 4]`). Given such tools in an interactive environment makes it easy to create vectors with integer values and see what happens when you use them as matrix indices in assignment statements.

Now set all the pixels whose values are greater than 0.5 to zero without resorting to loops. It is possible to do this in one statement. The `find()` function is likely your best bet here. Create a new figure and display the result (\$). You should understand how this works.

*Hint: An expression like `(im>0.5)` evaluates to a Boolean matrix, which is `true` where the condition holds. The `find` function returns a vector containing the indices of the true values. This vector can be used to index the image. You might ask how does this work when the matrix is 2D and indexes are arranged as 1D? Matlab lets you treat matrices as 1D vectors too, linearizing the matrix in **column-major** order.*

*More on **column-major** order. Think for a moment how 2D arrays are stored in memory. There are a number of options. One options is that the array is stored as a linear sequence of numbers, i.e., a 1D vector. This still leaves two alternatives. One is that the first row is followed by the second row (row-major), which is how C/C++ handle fixed arrays. The second is that the first column is stored in order, followed by the second column, and so on. This is **column-major** order which is used by Fortran and inherited by Matlab. This is an important issue to understand for those that might want to call Fortran routines from C which is a useful skill.*

Note: It is better to avoid loops where possible as Matlab is usually used as an interpreted language (you can compile it, but at this point I think you should consider using C/C++). If we use the built in matrix manipulations and functions to do the expensive work, Matlab can be reasonably fast.

More tricks for those that are interested (no deliverables)

Matlab makes it easy to write “vectorized” expressions without having to write for loops or if statements. For example, this will add all the values of the image:

```
sum(im(:))
```

The following will count the number of values greater than 0.9:

```
numel(find(im>0.9))
```

So will this:

```
sum(sum(im>0.9))
```

This will halve only those values greater than 0.9 (note the use of the `.` operator to do element-size multiplication of matrices):*

```
im = im - 0.5*im.*(im>0.9);
```

And so will this:

```
mask = (im>0.9);  
im = im.*~mask + im.*mask*0.5;
```

This will set 100 unique random pixels to zero:

```
p = randperm(numel(im));  
im(p(1:100)) = 0;
```

See `help elmat` for a list of interesting matrix manipulation and creation routines.

17. PCA (*)

Principle components analysis (PCA) is a very common method for transforming data for a variety of purposes. There is a reasonable chance that you have already been exposed to PCA. This problem and the next take a brief look behind the scenes. In computer vision PCA has been used as a way to represent data sets so that the “important” parts are exposed (and the parts you want to ignore, such as noise, are excluded). Representing images in this way enables pattern matching on the important parts. The first and perhaps most famous paper is “Face recognition using eigenfaces” by Turk and Pentland, 1991. Having said that, that famous paper was a bit flawed, and this simple technique runs out of steam quickly, so publishing such a paper would be hard today. Nonetheless, vision students should be familiar with PCA, as it is often used as part of vision systems.

The data matrix `Data/pca.txt` (also linked here <http://kobus.ca/teaching/cs477/data/pca.txt>) has two columns for X values and Y values, respectively. Read them in, and provide a plot of X vs Y (\$). Compute the covariance matrix for the data and report it (\$). Now consider changing the coordinate system by shifting the data and/or rotating it about the origin. Based on your plot, describe a more natural XY (Euclidean) coordinate system for this data (\$). The new coordinate system is like having an XY axis that is a shifted and rotated version of the standard one. Think a minute whether there are any reasonable options for “natural”, but if you are still confused, see the hint below. Under your new coordinate system, describe what you would expect the covariance matrix to be like (provide an “eyeball” estimate for the elements, and explain) (\$).

Hint: It is sometimes natural to “mean-center” the data which shifts data so that the mean is the origin. Second, if you believe the data is really one-dimensional, then you might want to make it so that one dimension is the true dimension, and the other gets all the noise.

18. PCA II (*)

To map the data to a space which might be the one you specified in the previous question, first zero center the data so that the mean of both X and Y is zero. Second rotate the data by applying the appropriate orthogonal transformation by multiplying all the data points by the appropriate rotation matrix. In particular, the matrix will have the eigenvectors of the covariance matrix as columns. See the function `eig()` to find that matrix. Provide the matrix and the computations which verify that it is orthogonal (\$). Now transform the data by using the matrix of eigenvectors (see tips below) and provide a plot of the resulting data. Make sure you use the same scale for the new X and Y axes (\$).

Note that it is easy to get the appropriate matrix transformation and its transpose mixed up, so if you are having trouble, try it the other way.

Now recompute the covariance matrix for the transformed data and provide the matrix (\$). Does it make sense given your previous answers? Explain a bit (\$). Finally, sum the two variance values and compare that sum with sum of the two variance values before the transformation (\$).

Tips: Here we want the new X to be the eigenvector associated with the biggest eigenvalue. Matlab usually[†] gives them to you smallest to largest. Also, we are using the orthogonal matrix to transform column

vectors, but you may have your data as rows. You can either deal with this head on, or consider using the transpose in the appropriate way.

[†]I say “usually” because it depends on the algorithm used. As of writing, the algorithm used usually gives this order, but it is possible that this is not exactly true. So, if you are relying on this, you want to check. The fact that it is possible that the default algorithm gives them out of order led to a confusing bug a few years ago in CS477/577.

What to Hand In

Hand in a Matlab program hw1.m and the PDF file hw1.pdf with the story of your efforts into D2L. Note that some of the exercises were just things that you should try---there is no corresponding code to hand in.